
SGBD

Jose Javier Alva Cornejo

Estructura de disco

≡ Crear Disco ≡

Nombre de disco a crear: disco

Ingrese número de superficies: 6

Ingrese número de pistas por superficie: 20

Ingrese número de sectores por pista: 20

Ingrese número de bytes por sector: 512

Ingrese número de sectores por bloque: 4

Disco creado exitosamente

ENTER para regresar ...

```
disco
├── superficie 0
├── superficie 1
├── superficie 2
├── superficie 3
├── superficie 4
├── superficie 5
└── free_space_bitmap
```

```
pista 3
pista 4
├── sector 0
├── sector 1
├── sector 10
├── sector 11
├── sector 12
├── sector 13
├── sector 14
```

```
// Se crea toda la estructura static library(disk_manager)
for (int s = 0; s < SURFACES; ++s) {
    std::string carpetaSuperficie = disk_name_used + "/superficie " + std::to_string(s);
    fs::create_directory(carpetaSuperficie);

    for (int p = 0; p < TRACKS_PER_SURFACE; ++p) {
        std::string carpetaPista = carpetaSuperficie + "/pista " + std::to_string(p);
        fs::create_directory(carpetaPista);

        for (int sec = 0; sec < SECTORS_PER_TRACK; ++sec) {
            std::string archivoSector = carpetaPista + "/sector " + std::to_string(sec);

            if (bin) {
                std::ofstream out(archivoSector, std::ios::binary | std::ios::trunc);
                std::vector<char> zeros(SECTOR_SIZE, 0);
                out.write(zeros.data(), SECTOR_SIZE);
                out.close();
            } else {
                std::ofstream(archivoSector + ".txt").close();
            }
        }
    }
}
```

Sector 0



sector 0.txt

```
Superficies: 6, tracks: 20, sectores: 20,  
tamano de sector: 512,sectores por  
bloque: 4--
```



```
/*  
 * Representa metadata esencial de un disco  
 */  
#pragma pack(push, 1)  
struct Sector_0 {  
    uint16_t surfaces{};  
    uint16_t tracks_per_surf{};  
    uint16_t sectors_per_track{};  
    uint16_t sector_size{};  
    uint16_t sectors_per_block{};  
};  
#pragma pack(pop)
```

Sector 1



sector 1.txt

0titanic

Max_reg_size: 216

Fixed_regs: 1

First page id: 1

Last page id: 99

n_columns 12

PassengerId 2 4-

Survived 2 4-

Pclass 2 4-

Name 6 100-

Sex 6 20-

Age 5 8-

SibSp 2 4-

Parch 2 4-

Ticket 6 20-

Fare 5 8-

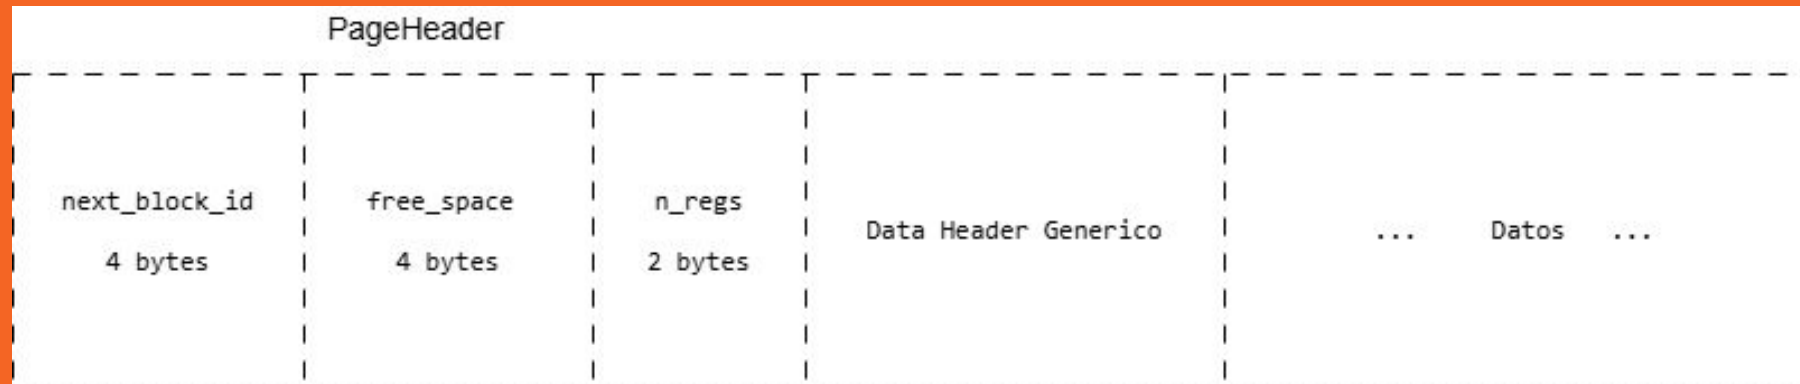
Cabin 6 20-

Embarked 6 20-

```
/*
 * Representa las ubicaciones de todas las relaciones
 * en disco(sector reservado 1)
 * A cada tabla le corresponde un bloque para
 * guardado de metadata
 */
struct Sector_1 {
    uint8_t n_tables{};
    std::vector<uint32_t> table_block_ids{};
};
```

Estructura de paginas

PageHeader



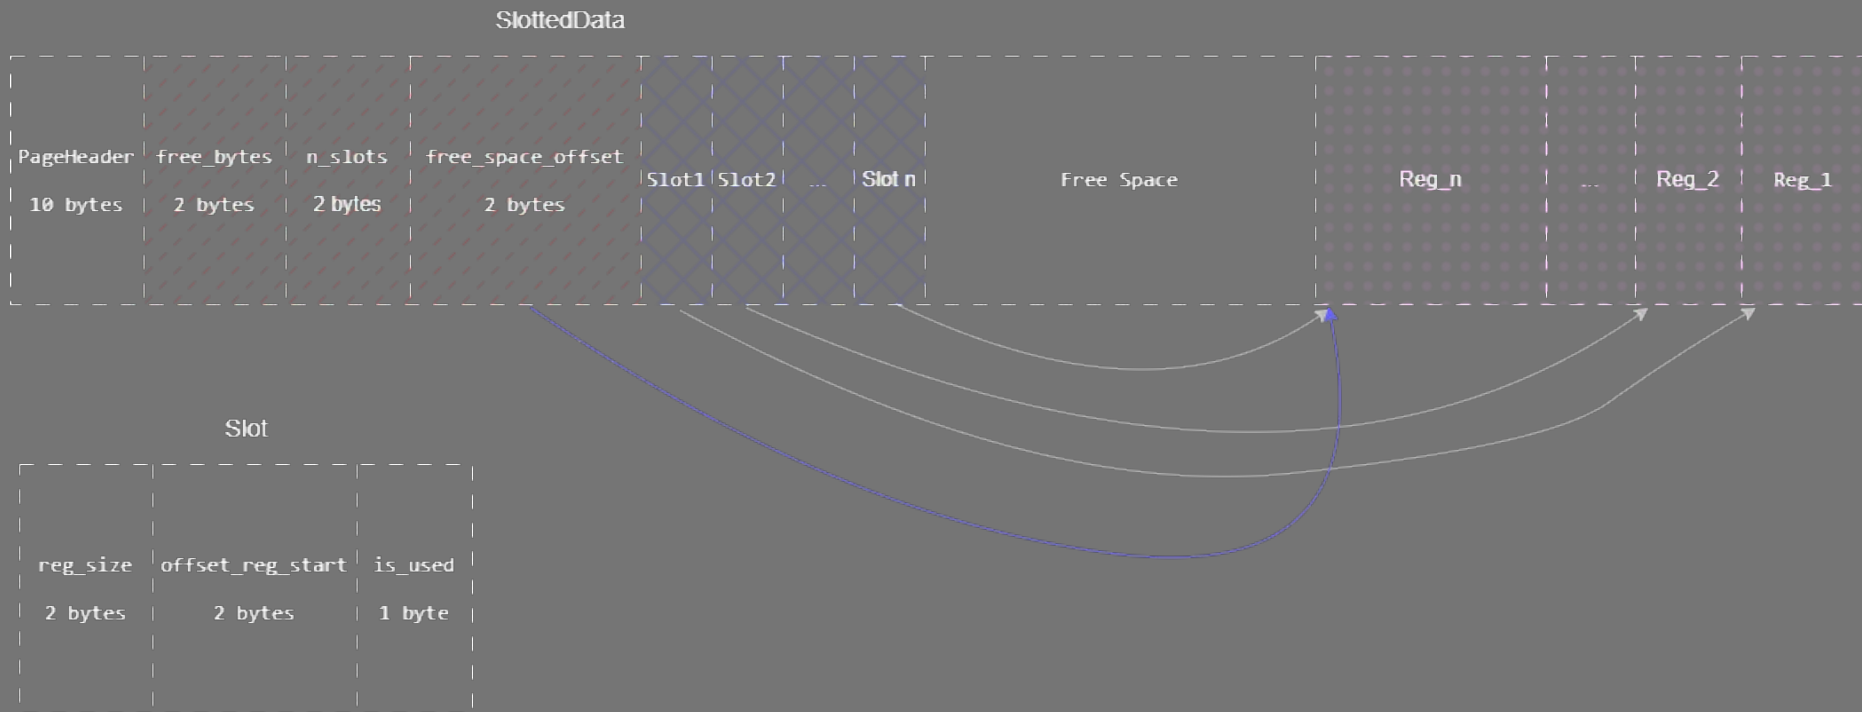
```
#pragma pack(push, 1)
struct PageHeader {
    uint32_t next_block_id{};
    uint32_t free_space{};
    uint16_t n_regs{};
};
#pragma pack(pop)
```

FixedDataHeader



```
/*
 * @struct Estructura para metadata de bloques interno de file donde
 * registros son de size fijos
 * @note El bitset tiene tamaño de la max cantidad de registros guardables
 */
struct FixedDataHeader {
    // Respecto a bloque completo, se resta tamaño mismo de header
    uint32_t free_bytes{};
    uint32_t reg_size{};
    uint16_t max_n_regs{};
    boost::dynamic_bitset<unsigned char>
        free_register_bitmap;
};
```

SlottedDataHeader



Ejemplo de serialización

```
template <typename Iter>
inline void serialize_fixed_block_header(const FixedDataHeader &header, Iter &out_it) {
    write_v(out_it, header.free_bytes);
    write_v(out_it, header.reg_size);
    write_v(out_it, header.max_n_regs);

    size_t bitset_byte_size = (header.max_n_regs + CHAR_BIT - 1) / CHAR_BIT;

    boost::to_block_range(b: header.free_register_bitmap, result: out_it);
    std::advance(out_it, bitset_byte_size);
}
```

Tabla
First_p
age_id

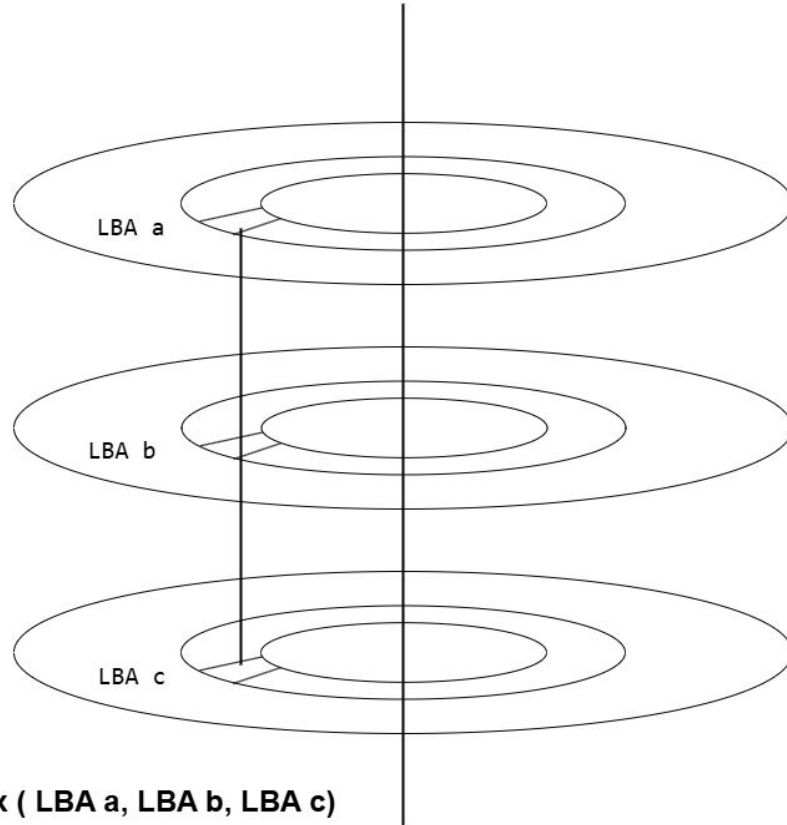
Pagina/Bloque		PageHeader		
		free_space	n_regs	next_block_id
		GenericDataHeader		
		...Registros...		
Sector 1				
Sector 2		...Registros...		
Sector 3		...Registros...		

Pagina/Bloque		PageHeader		
		free_space	n_regs	next_block_id
		GenericDataHeader		
		...Registros...		
Sector 1				
Sector 2		...Registros...		
Sector 3		...Registros...		

NULL

Bloques

Diccionario de datos/FreeBlockMap



Asignacion vertical

Bloque agrupa ciertos
sectores

```
struct FreeBlockMap {
    std::vector<std::pair<bool, std::vector<uint32_t>>> blocks{};

    bool is_block_free(uint32_t block_id) {
        if (block_id ≥ blocks.size())
            throw std::runtime_error("block_id fuera de rango");

        return !blocks[block_id].first;
    }

    uint32_t get_ith_lba(uint32_t block_id, size_t ith_lba) {
        if (block_id ≥ blocks.size())
            throw std::runtime_error("block_id fuera de rango");

        return blocks[block_id].second[ith_lba];
    }

    void set_block_used(uint32_t block_id) {
        if (block_id ≥ blocks.size())
            throw std::runtime_error("block_id fuera de rango");

        blocks[block_id].first = true;
    }
};
```

Next_page_id: 2 N_registers: 18 Free_space/capacity: 4/404

free_space_offset: 110 # Slots: 18

|1|99|1901|

|1|128|1773|

|1|107|1666|

|1|119|1547|

|1|97|1450|

|1|89|1361|

Superficie: 3 Track: 0 Sector: 1

103"Braund; Mr. Owen Harris"male22.00000010A/5 211717.250000 S
211"Cumings; Mrs. John Bradley (Florence Briggs Thayer)"female38.00000010PC 1759971.
283300C85 C
313"Heikkinen; Miss. Laina"female26.00000000STON/O2. 31012827.925000 S
411"Futrelle; Mrs. Jacques Heath (Lily May Peel)"female35.0000001011380353.
100000C123 S

Superficie: 2 Track: 0 Sector: 1

503"Allen; Mr. William Henry"male35.000000003734508.050000 S
603"Moran; Mr. James"male0.000000003308778.458300 Q
701"McCarthy; Mr. Timothy J"male54.000000001746351.862500E46 S
803"Palsson; Master. Gosta Leonard"male2.0000003134990921.075000 S
913"Johnson; Mrs. Oscar W (Elisabeth Vilhelmina Berg)"female27.0000000234774211.
133300 S

Superficie: 1 Track: 0 Sector: 1

1012"Nasser; Mrs. Nicholas (Adele Achem)"female14.0000001023773630.070800

Manejo de Registros

Busqueda
Tabla en
sector 1



Se encuentra,
se empieza con
first_page_id



Se itera toda
pagina, se busca
un Page_header
con espacio



Se llega a
NULL, crea
nueva pagina



Se tiene una
pagina
insertable

Inserción de un registro

Next_page_id: 600 N_registers: 1 Free_space/capacity: 1762/488
Register size: 216 Max registers: 9 Free_register_bitmap: 000000001

Superficie: 0 Track: 0 Sector: 11

103Braund; Mr. Owen

Harris

male 22.00000010A/5 21171 7.250000

S_____

Superficie: 1 Track: 0 Sector: 11

Linea “wrappeada” por motivos de visualizar, en realidad es una linea continua

titanic.csv ×

☰ bloque 16.txt ×

☰ sector 11.txt ×

103Braund; Mr. Owen
Harris
male22.00000010A/5 211717.250000
S_____

Inserción hasta exceder un sector

Next_page_id: 600 N_registers: 3 Free_space/capacity: 1330/488
Register size: 216 Max registers: 9 Free_register_bitmap: 000000111

Superficie: 0 Track: 0 Sector: 11

103Braund; Mr. Owen Harris

male

22.00000010A/5 21171

7.250000

S

211Cumings; Mrs. John Bradley (Florence Briggs Thayer)

female

38.00000010PC 17599

71.283300C85

C

Superficie: 1 Track: 0 Sector: 11

313Heikkinen; Miss. Laina

female

26.00000000STON/02. 3101282

7.925000

S

Next_page_id: 600 N_registers: 3 Free_space/capacity: 1330/488

Register size: 216 Max registers: 9 Free_register_bitmap: 000000111

Superficie: 0 Track: 0 Sector: 11

103Braund; Mr. Owen

Harris

male

22.00000010A/5 21171

7.

250000

S

211Cumings; Mrs. John Bradley (Florence Briggs
Thayer)

female

38.00000010PC 17599

71.

283300C85

C

Superficie: 1 Track: 0 Sector: 11

313Heikkinen; Miss.

Laina

female

26.00000000STON/02. 3101282

7.

925000

S

titanic.csv x ≡ bloque 16.txt x ≡ **sector 11.txt** x

103Braund; Mr. Owen
Harris
male 22.00000010A/5 21171 7.250000
S_____

211Cumings; Mrs. John Bradley (Florence Briggs
Thayer) female 38.00000010PC
17599 71.283300C85 C_____

titanic.csv x ≡ bloque 16.txt x ≡ **sector 11.txt** x

313Heikkinen; Miss.
Laina
female 26.000000000STON/02. 3101282 7.925000
S_____

Insertión de un registro variable

```
Next_page_id: 600 N_registers: 1 Free_space/capacity: 1851/489
free_space_offset: 1872 # Slots: 1
|1|128|1872|
|
Superficie: 3 Track: 0 Sector: 1
1903Vander Planke; Mrs. Julius (Emelia Maria Vandemoortele)female31.0000001034576318.
000000S
```

```
titanic.csv × | ≡ bloque 1.txt × | ≡ sector 1.txt ×
1903Vander Planke; Mrs. Julius (Emelia Maria Vandemoortele)female31.0000001034576318.
000000S
```

Insertión variable hasta exceder sector

```
titanic.csv × | ≡ bloque 1.txt ×
Next_page_id: 600 N_registers: 5 Free_space/capacity: 1456/469
free_space_offset: 1497 # Slots: 5
|1|128|1872|
|1|94|1778|
|1|91|1687|
|1|92|1595|
|1|98|1497|
|
Superficie: 3 Track: 0 Sector: 1
1903Vander Planke; Mrs. Julius (Emelia Maria Vandemoortele)female31.0000001034576318.
000000S
2013Masselmani; Mrs. Fatimafemale0.0000000026497.225000C
2102Fynney; Mr. Joseph Jmale35.0000000023986526.000000S
2212Beesley; Mr. Lawrencemale34.0000000024869813.000000D56S

Superficie: 2 Track: 0 Sector: 1
2313McGowan; Miss. Anna Anniefemale15.000000003309238.029200Q
```

titanic.csv x | ≡ bloque 1.txt x | ≡ **sector 1.txt** x

1903	<u>Vander Planke</u> ; Mrs. Julius (<u>Emelia Maria Vandemoortele</u>)	<u>female</u>	<u>31.000000</u>	<u>1034576318.000000</u>		
		S				
2013	<u>Masselmani</u> ; Mrs. <u>Fatima</u>	<u>female</u>	<u>0.000000</u>	<u>0026497.225000</u>		C
2102	<u>Fynney</u> ; Mr. Joseph <u>Jmale</u>	<u>35.000000</u>	<u>0023986526.000000</u>			S
2212	<u>Beesley</u> ; Mr. <u>Lawrence</u>	<u>male</u>	<u>34.000000</u>	<u>0024869813.000000</u>	<u>D56</u>	S

titanic.csv x | ≡ bloque 1.txt x | ≡ **sector 1.txt** x

2313	<u>McGowan</u> ; Miss. Anna <u>Annie</u>	<u>female</u>	<u>15.000000</u>	<u>003309238.029200</u>		Q
------	--	---------------	------------------	-------------------------	--	---

Insert Fixed implementación

```
// Serializamos todo el registro static library(megatron)
auto register_bytes: std::vector<unsigned char> = serialize_register(table_metadata,
|                                     &: values);

// Se busca pagina a insertar
std::vector<unsigned char> page;
uint32_t insert_page_id;

insert_page_id = get_insertable_page(&: page,
|                                     block_id: table_metadata.first_page_id,
|                                     reg_size: table_metadata.max_reg_size);

// Paginas sin espacio suficiente
if (insert_page_id == disk.NULL_BLOCK)
|   insert_page_id = add_new_page_to_table(&: table_metadata);

// Calculamos posicion donde insertar
size_t free_reg_pos = serial::find_free_reg_pos(header: fixed_data_header);
byte_offset_free_reg = serial::calculate_reg_offset(header: fixed_data_header,
|                                                    nth_reg: free_reg_pos);

disk.write_block(&block_bytes: insert_page_bytes, block_id: insert_page_id);
```


Insert Variable Implementación

[illegible]

Select Fixed

```
// Se saca metadata relevante static library(megatron)
serial::PageHeader page_header;
serial::FixedDataHeader fixed_data_header;
serial::SlottedDataHeader slotted_data_header;

page_header = serial::deserialize_page_header(page_bytes_it);

fixed_data_header = serial::deserialize_fixed_data_header(page_bytes_it);
for (size_t i{}; i < fixed_data_header.max_n_regs; ++i) {
    if (fixed_data_header.free_register_bitmap.at(i)) { // Registro existe
        auto register_bytes = get_ith_register_bytes(table_metadata, page_header, fixed_data_header, i);
        auto register_values = deserialize_register(table_metadata, register_bytes);

        // Si hay condicion
        if (col_index < table_metadata.n_cols && register_values[col_index] != cond_val)
            continue;

        for (auto &v : register_values)
            std::cout << SQL_type_to_string(v) << " | ";

        std::cout << std::endl;
    }
}
```


Select Slotted

```
// Se saca metadata relevante static library(megatron)
serial::PageHeader page_header;
serial::FixedDataHeader fixed_data_header;
serial::SlottedDataHeader slotted_data_header;

page_header = serial::deserialize_page_header(page_bytes_it);

slotted_data_header = serial::deserialize_slotted_data_header(page_bytes_it);
for (size_t i{}; i < slotted_data_header.n_slots; ++i) {
    if (slotted_data_header.slots[i].is_used) { // Registro existe
        auto register_bytes = get_ith_register_bytes(table_metadata, page_header, -
                                                    slotted_data_header, page_bytes, i);
        auto register_values = deserialize_register(table_metadata, register_bytes);

        // Si hay condicion
        if (col_index < table_metadata.n_cols && register_values[col_index] != cond_val)
            continue;

        for (auto &v : register_values)
            std::cout << SQL_type_to_string(v) << " | ";

        std::cout << std::endl;
    }
}
```

Eliminar fijo -> Similar a select

```
if (register_values[col_index] ≠ cond_val)
|   continue;
|
// Si cumple condicion, delete
// Solo marcamos como libre, ya que todo es fijo se reescribira luego
fixed_data_header.free_bytes += fixed_data_header.reg_size;
fixed_data_header.free_register_bitmap[i] = false;
page_header.free_space += fixed_data_header.reg_size;
page_header.n_regs--;
```

Eliminar variable

```
if (register_values[col_index] ≠ cond_val)
|   continue;
|
// Cumple condicion, solo marcamos slot como disponible
// No se actualiza free_space, este depende de un compactar
slotted_data_header.slots[i].is_used = false;
page_header.n_regs--;
```

Luego de delete PClass 3

```
Next_page_id: 2 N_registers: 6 Free_space/capacity: 4/404  
free_space_offset: 110 # Slots: 18  
|0|99|1901|  
|1|128|1773|  
|0|107|1666|  
|1|119|1547|  
|0|97|1450|  
|0|89|1361|  
|1|95|1266|  
|0|103|1163|  
|0|124|1039|  
|1|110|929|  
|0|107|822|  
|1|99|723|  
|0|106|617|  
|0|100|517|  
|0|111|406|  
|1|107|299|  
|0|93|206|  
|0|96|110|  
|
```

Superficie: 3 Track: 0 Sector: 1

211"Cumings; Mrs. John Bradley (Florence Briggs Thayer)"female38.00000010PC 1759971.
283300C85 C

411"Futrelle; Mrs. Jacques Heath (Lily May Peel)"female35.0000001011380353.
100000C123 S

Superficie: 2 Track: 0 Sector: 1

Referencias

1. Silberschatz, A., Korth, H. F., & Sudarshan, S. (2011). Database system concepts (6th ed.). McGraw-Hill.
 2. Garcia-Molina, H., Ullman, J. D., & Widom, J. (2008). Database systems: The complete book (2nd ed.). Pearson Education.
 3. Universiteit Antwerpen. (s.f.). Architecture of DBMS, Organizational matters & Recap [Presentación de clase]. Universiteit Antwerpen.
-